

FACOLTÀ DI INGEGNERIA DELL'INFORMAZIONE, INFORMATICA E STATISTICA

Corso di Laurea Magistrale in Computer Science

Thesis Defense

Robustness Of Deep Neural Networks Using Trainable Activation Functions

Candidate:


Federico Peconi

Thesis Advisor:

Prof. Simone Scardapane

Academic Year:


2019/2020



SAPIENZA
UNIVERSITÀ DI ROMA


Intriguing Properties Of Neural Networks

- ❑ Neural Networks are brittle w.r.t. *non-random*, small perturbations of the input [Christian Szegedy, Wojciech Zaremba et al. 2014]



- ❑ *Adversarial Examples* can be found in several different domains other than image classification:

1. Natural Language Processing [Moustafa Alzantot et al. 2018]
 2. Automatic Speech Recognition [Yao Qin et al. 2019]
- and so on ...



[Yuan Gong, Christian P. 2018]

- ❑ Can be computed through *backpropagation* and *gradient ascent*, optimizing:

$$\max_{\delta \in \Delta} \ell(f_{\theta}(x + \delta), y)$$

Agenda




- Motivations and Current Effective(?) Defenses
- Kernel-Based Activation Functions (KAFs)
- The role of activation functions in Adversarial Training
- VGG architectures-inspired results
- ResNet20 results
- Contributions and Future Works

Motivations

❑ High Security Concerns:

- *Safety-critical* applications not yet reliable
- Malware, intrusion detection
- Face recognition

...



[Mahmood Sharif, Sruti Bhagavatula et al. 2016]



[Kevin Eykholt et al. 2018]


❑ We want to develop human-aligned models:

- Basic idea: humans do not suffer from such brittleness
- *Robustness* and *Interpretability* seem to be strictly connected
- “Adversarial examples are not bugs, they are features” [Madry et al. 2019]


primate



std saliency



robust saliency

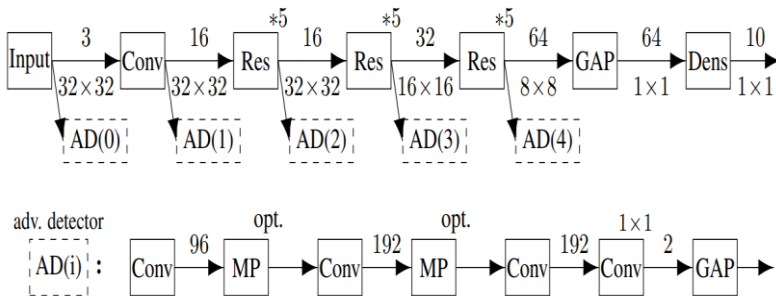


[Dimitris Tsipras, Shibani Santurkar et al. 2019]

Defenses

Detection Methods

- ❑ Add a “patch” to the network which is in turn another neural network.
- ❑ Such NN, is used as a *detector* to spot adversarial perturbations from the analysis of layer’s *statistics*.
- ❑ Training samples extended with a *binary label* denoting the presence of an adversary
- 😊 Easy to perform, tractable computational overhead
- ☹️ Tend to *overfit* on the specific attack used during training



[H. Metzen et al. 2017]

Adversarial Training (AT)

- ❑ Train to approximate the *min*, *max*:

$$\min_{\theta} \frac{1}{|S|} \sum_{x,y \in S} \max_{\delta \in \Delta} \ell(f_{\theta}(x + \delta), y)$$

- ❑ Inner maximization is computed crafting adversarial examples on-the-fly
- ❑ *Running-time* boosted with novel techniques: [E. Wong, L. Rice 2020]
- 😊 Only *unbroken* known defense, believed to be *truly robust*
- ☹️ Robustness accuracy up to 50%
Std Accuracy ~ 80 %
(CIFAR10, delta <= 8/255)

Provable Robustness

- ❑ Aims to give a *formal certification* of the robustness of a model.
- ❑ We want to guarantee that, for any allowed perturbation our model will always predict the correct class
- ❑ Problem approached through different techniques:
 - Lipschitz Regularization
 - Linear Programming
 - Semi-Definite Programming
 - Randomized Smoothing
- 😊 Provides provably robust models (up to computationally negl. probability)
- ☹️ Currently intractable for large-scale networks, architecture-specific implementations

Kernel-Based Activation Functions

□ Along with ‘traditional’ *fixed* activation functions (e.g. ReLU, ELU, Tanh, Swish, ..), we can devise *adaptive* activation functions whose shape is learned through the optimization of parameters present in their formulation:

- Adaptive Piece-Wise Linear Activation Functions (APLs) [F. Agostinelli 2014]
- Spline Activation Functions (SAFs) [S. Scardapane, M. Scarpiniti 2017]
- Maxout layers [I. Goodfellow et al. 2013]

□ In [S. Scardapane, S. V. Vaerenbergh 2017], authors introduce a novel class of trainable activation functions called *kernel-based activation functions (KAFs)*:

$$\text{KAF}(x) = \sum_{i=1}^D \alpha_i \kappa(x, d_i)$$

- Where $\{\alpha_i\}_1^D$ are the weights to train, called *mixing coefficients*.
- $\kappa : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ is a 1-dimensional kernel method.
- $\{d_i\}_1^D$ are the kernel’s dictionary elements and D the dictionary size, which is user-defined.
- First derivatives: $\frac{\partial \text{KAF}(x)}{\partial \alpha_i} = \kappa(x, d_i), \quad \frac{\partial \text{KAF}(x)}{\partial x} = \sum_{i=1}^D \alpha_i \frac{\partial \kappa(x, d_i)}{\partial x}$

KAF Design Details

□ Dictionary elements uniformly sampled around 0 with step size Δ

□ Gaussian kernel: $k(x, d_i) = \exp\{-\gamma(x - d_i)^2\}$

- Heuristic $\gamma := \frac{1}{6\Delta^2}$

- Guarantees a *local* effect of the mixing coefficients w.r.t. the overall shape \Rightarrow ease optimization

□ Mixing coefficients are *initialized* either:

- *Randomly* from a normal distribution



- With *ridge-regression* if we want to *approximate* any fixed activation function $\sigma: \mathbb{R} \rightarrow \mathbb{R}$

$$\sigma(\mathbf{d}) = \mathbf{t} \implies \alpha = (\mathbf{K} - \varepsilon \mathbf{I})^{-1} \mathbf{t}$$

with: $K_{i,j} = \kappa(d_i, d_j)$

□ Such constructed KAFs satisfy:

- *Universal* approximability
- *Smoothness* over the entire domain
- Only *linear number of weights* introduced in the model
- Mixing coefficients admit common *regularization techniques*



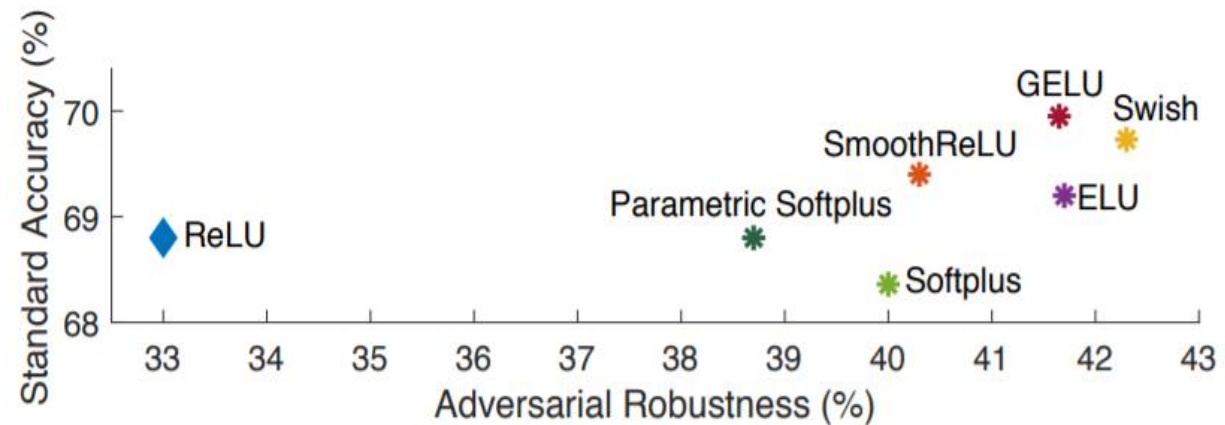
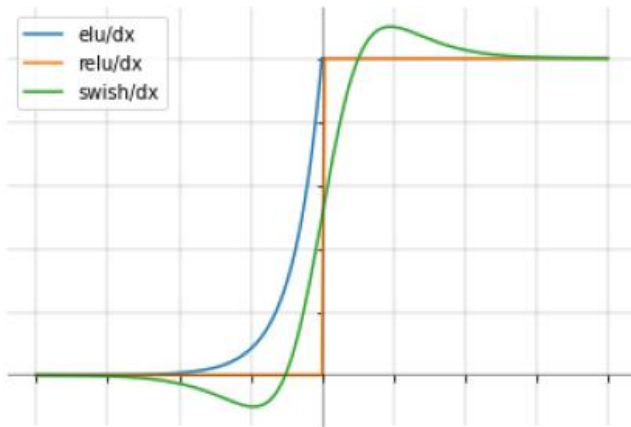
Smooth Activation Functions and AT

Can activation functions play a role in making neural networks more robust?

- Remark: to perform adversarial training we compute *twice* the number of *gradient computations* used for standard training, that is, both for the *inner maximization* and the *outer minimization*.

$$\min_{\theta} \frac{1}{|S|} \sum_{x,y \in S} \max_{\delta \in \Delta} \ell(f_{\theta}(x + \delta), y)$$

- In [Chiang Xie, M. Tan, B. Gong et al, 2020] authors conjectured that the widely-used ReLUs activation functions *weaken* AT due to their *non-smooth* nature
- Empirical results: by *enforcing smoothness*, replacing ReLUs with smooth counterparts, it is possible to train the network against stronger adversarial examples, which eventually leads to a *significant improvement* in the robustness



Hypothesis

```
1: procedure FBF_AT( $epochs = T, dataset\_size = M, \epsilon$ )
2:   for  $t \leftarrow 1, T$  do                                     ▷ For each epoch
3:     for  $i \leftarrow 1, M$  do                                   ▷ For each minibatch
4:        $\delta \leftarrow \mathcal{U}(-\epsilon, \epsilon)$            ▷ Random init
5:        $\delta \leftarrow \delta + 1.25\epsilon \cdot \text{sign}(\nabla_{\delta} \text{LS}(f_{\theta}(x + \delta), y))$    ▷ FGSM
6:        $\delta \leftarrow \mathcal{P}(\delta)$ 
7:        $\theta = \theta - \nabla_{\theta} \ell(f_{\theta}(x_i + \delta), y_i)$    ▷ Update model weights
8:     end for
9:   end for
10: end procedure
```

Smoothness (green arrow pointing to line 5)

Flexibility of KAFs (green arrow pointing to line 8)

□ KAFs might be able to enhance even more the robustness and the accuracy of adversarially-trained models.

- By leveraging *flexibility*, KAFs proved to outperform fixed activation functions in a variety of (adv. free) tasks
- Using the Gaussian kernel we can enforce *smoothness*
- Formally this translates to: potential to improve *both outer and inner optimization*, respectively

Experiments Set-up



❑ Training:

- Fast Adversarial Training [E. Wong, L. Rice 2020] until convergence
- *One-Cycle* learning rate policy [Leslie N. Smith 2018] and *mixed-integer precision*
- *CIFAR10* image classification dataset
- *SGD* optimizer and *batch size* of 128
- Categorical *cross-entropy* loss
- *Dictionary size* for KAFs is $D = 20$

❑ Adversarial examples computed through *Projected-Gradient Descent (PGD)*

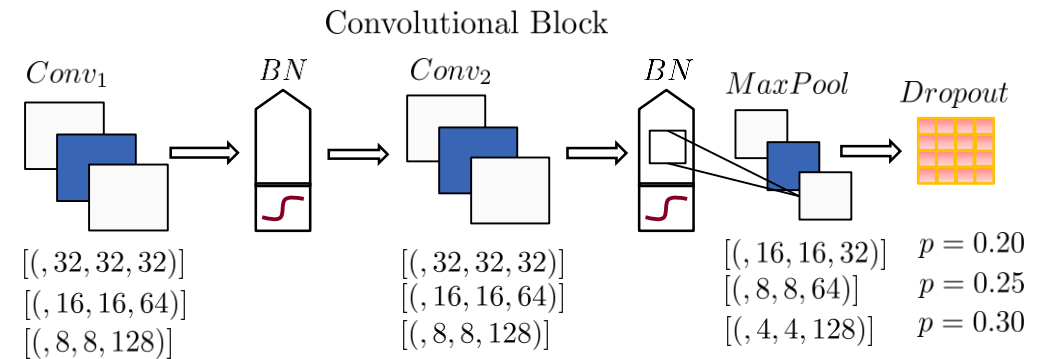
- Any perturbation lies in the L_∞ norm-ball bounded by $\epsilon = 8/255$
- 50 iterations sufficient for convergence
- 10 random restarts

❑ Evaluation

- *Robustness* i.e. accuracy of the model over PGD-perturbed test set
- *Accuracy* i.e. standard accuracy over clean test set

VGG Results


- ❑ VGG-inspired **architecture**: *three consecutive convolutional blocks*
- ❑ Followed by: *flatten layer, fully-connected layer w/ 128 units, dropout and a decision layer with softmax.*



- ❑ Train and evaluate seven different architectures that differ *only* on the activation function employed:

| Activation | Accuracy | Robustness |
|------------|---------------|---------------|
| ReLU | 78.1% | 45.41% |
| ELU | 73.99% | 45.49% |
| KAF | <u>77.93%</u> | <u>46.8%</u> |
| KAF_ELU | 74.93% | 45.4% |
| Swish | 76.99% | 47.73% |
| KAF_Swish | 78.85% | 46.33% |
| KAF_Free | 77.54% | 46.2% |


Tab1: Bold denotes best result, underline best trade-off



ResNet20 Results

- ❑ *Max-pooling* layers are *non-smooth*, which makes our hypothesis misplaced in the first place. We need to enforce smoothness throughout the whole model to perform a meaningful evaluation.
- ❑ Due to its relatively simple implementation and strong performances we decided to repeat the experiment using a *residual neural network* (ResNet).
- ❑ We experienced an *exploding gradient* issue using KAFs in deep (>50 layers) scenarios, thus we chose a ResNet20 architecture as described in the original paper [Kaiming He, X. Zhang et al. 2015]
- ❑ *Data augmentation*: 4 pixels are 0-padded on each side of the image and a 32×32 crop is randomly sampled from the resulting image or its horizontal flip.

| Activation | Accuracy | Robustness |
|------------|---------------|---------------|
| ReLU | 77.52% | 50.17% |
| ELU | 75.42% | 51.69% |
| KAF | 78.08% | 52.23% |
| KAF_ReLU | 79.65% | 51.44% |
| KAF_ELU | 78.69% | 52.76% |
| Swish | <u>78.68%</u> | 53.88% |
| KAF_Swish | 79.21% | 52.10% |
| KAF_Free | 78.88% | 52.07% |



- ❑ Overall robustness improved wrt VGG case \implies ResNet better suited for adversarial training (approx. same number of weights)
- ❑ KAFs reach *best std accuracies*, nevertheless, *Swish* function allows again for *strongest robustness*

Conclusions and Contributions

Sub-optimality in terms of robustness might be explained with *the introduction of new parameters* which comes with *KAFs* and ultimately with an *increased non-linearity* of the model's landscape. Despite being leveraged by AT, it can also be exploited by any *gradient-based* attack.

Fix: Try to use *Madry's original adversarial training* procedure instead [A. Madry, A. Makelov, L. Schmidt et al. 2017]

- ❑ Presented results can also be seen as a *contribution of evidences to the smoothness thesis* [Chiang Xie, M. Tan, B. Gong et al, 2020]
- ❑ We gave a working implementation in TF2 of cutting-edge procedures such as [E. Wong, L. Rice 2020] currently lacking even in third- party libraries. Moreover, we extended the *KAF layer* implementation for Keras.
- ❑ **Future Works:**
 - Test KAFs robustness also in *provable* scenarios such as *Lipschitz-regularization training*. [Y. Tsuzuku, I. Sato et al. 2018]
 - *Scale* the proposed defense towards larger datasets such as *ImageNet* to see if results are coherent.
 - Assess robustness using more-sophisticated attacks: *C&W*, *DeepFool* or, more importantly, *adaptive attacks* [F. Tramer, N. Carlini et al. 2020]

QUESTIONS?

(Code and Adversarially-trained Models)

THANK YOU!